

GROUP 1: KARU PHASE 2 PROJECT

- Sheilla Macharia
- Margaret Mondia
- Benard Kariuki
- Ibrahim Mohamud
- Bravian Ashono
- Isaac Ndung'u

BUSINESS UNDERSTANDING

In the fiercely competitive entertainment sector, studios and production companies must have a thorough understanding of the elements that go into a successful film. Data-driven choices can greatly increase profitability and viewer engagement in the face of rising production costs and shifting audience preferences.

In order to find important information that can guide production and marketing strategies, this project will examine movie data from a number of reliable sources, such as TMDB, Box Office Mojo, IMDb, and Rotten Tomatoes. Finding the optimal movie runtime—the duration that optimises audience ratings and overall success—is one of the main goals.

The project aims to advise studios on the ideal length of content for engagement and financial performance by examining trends in movie runtimes, earnings, and ratings.

BUSINESS QUESTIONS

1. What trends exist across genres, release periods, and production budgets in relation to success?
2. What is the ideal runtime minutes for a movie?
3. Can early popularity forecast long-term success?

Description of the Datasets

Several datasets were combined to provide a comprehensive view of movie characteristics and performance:

Dataset	Description	Key Columns
tmdb.movies.csv	Contains movie details from The Movie Database (TMDB).	id, original_title, popularity, release_date, vote_average, vote_count
bom.movie_gross.csv	Box Office Mojo data on domestic and foreign grosses.	title, studio, domestic_gross, foreign_gross, year
tn.movie_budgets.csv	Budget and revenue data for movies.	id, release_date, production_budget, domestic_gross, worldwide_gross
clean_movie_basics.csv	Basic movie info including year and genre.	movie_id, primary_title, start_year, genres

Dataset	Description	Key Columns
rt.movie_i nfo.csv	Rotten Tomatoes metadata with movie ratings and runtime.	id, rating, genre, box_office, runtime, studio
<pre># Import all the relevant libraries import pandas as pd import numpy as np import seaborn as sns import matplotlib.pyplot as plt %matplotlib inline import sqlite3 from scipy import stats import warnings warnings.filterwarnings("ignore")</pre>		

IMDB DataBase

We will use sql to query data from the zipped data and connect the relevant tables

```
conn = sqlite3.connect("C:\\\\Users\\\\Sheilla Macharia\\\\Documents\\\\MoringaProjects\\\\assignments\\\\Phase_2_Project\\\\data\\\\im.db")

#Use sql to query the data now
movie_basics_ratings = pd.read_sql("""
SELECT *
FROM movie_basics
JOIN movie_ratings
USING (movie_id) ;"""
, conn)

#Check the first five entries of our data
movie_basics_ratings.head()

      movie_id           primary_title
original_title \
0  tt0063540           Sunghursh
Sunghursh
1  tt0066787  One Day Before the Rainy Season
Din
2  tt0069049      The Other Side of the Wind
The Other Side of the
Wind
3  tt0069204           Sabse Bada Sukh
Sabse Bada
Sukh
4  tt0100275      The Wandering Soap Opera
Errante
                                         genres  averagerating
numvotes
0            2013        175.0  Action,Crime,Drama
77
77
```

1	2019	114.0	Biography, Drama	7.2
43				
2	2018	122.0	Drama	6.9
4517				
3	2018	NaN	Comedy, Drama	6.1
13				
4	2017	80.0	Comedy, Drama, Fantasy	6.5
119				

The dataframe above has 8 columns:

1. `movie_id` - Unique Identifier of each movie entry
2. `primary_title` - The primary title of the movie
3. `original_title` - The Actual title of the movie
4. `start_year` - The year the movie was premiered

Box Office Mojo

```
#Load the dataset and read the first 5 rows
movie_gross= pd.read_csv("data/bom.movie_gross.csv")
movie_gross.head()

          title studio domestic_gross
\0           Toy Story 3     BV  415000000.0
1           Alice in Wonderland (2010)  BV  334200000.0
2  Harry Potter and the Deathly Hallows Part 1  WB  296000000.0
3           Inception      WB  292600000.0
4           Shrek Forever After  P/DW  238700000.0

    foreign_gross  year
0       652000000  2010
1       691300000  2010
2       664300000  2010
3       535700000  2010
4       513900000  2010

# Identify the number of columns and rows
movie_gross.shape
(3387, 5)
```

The box office dataframe has 5 columns and 3,387 entries. Each column holds a different aspect of the dataset:

1. `title` Contains the title of the movies

2. `studio` column has abbreviated names of the studios producing the respective movie
3. `domestic_gross` column shows the amount of revenue generated by the movie through sales, locally
4. `foreign_gross` shows the amount of revenue generated by the movie internationally
5. `year` shows the year when the movie was premiered

The Movie Data Base

```
#Load the dataset and read the first 5 rows
TMDB = pd.read_csv("data/tmdb.movies.csv")
TMDB.head()

      Unnamed: 0      genre_ids      id original_language \
0            0      [12, 14, 10751]  12444                  en
1            1      [14, 12, 16, 10751]  10191                  en
2            2      [12, 28, 878]   10138                  en
3            3      [16, 35, 10751]     862                  en
4            4      [28, 878, 12]   27205                  en

                                         original_title  popularity
release_date \
0  Harry Potter and the Deathly Hallows: Part 1    33.533  2010-11-
19
1                      How to Train Your Dragon    28.734  2010-03-
26
2                          Iron Man 2    28.515  2010-05-
07
3                          Toy Story    28.005  1995-11-
22
4                         Inception    27.920  2010-07-
16

                                         title  vote_average
vote_count
0  Harry Potter and the Deathly Hallows: Part 1        7.7
10788
1                      How to Train Your Dragon        7.7
7610
2                          Iron Man 2        6.8
12368
3                          Toy Story        7.9
10174
4                         Inception        8.3
22186

df = pd.read_csv("data/tn.movie_budgets.csv")
df.head()

      id release_date                                movie \
0     1 Dec 18, 2009                                Avatar
```

```

1 2 May 20, 2011 Pirates of the Caribbean: On Stranger Tides
2 3 Jun 7, 2019 Dark Phoenix
3 4 May 1, 2015 Avengers: Age of Ultron
4 5 Dec 15, 2017 Star Wars Ep. VIII: The Last Jedi

  production_budget domestic_gross worldwide_gross
0      $425,000,000    $760,507,625   $2,776,345,279
1      $410,600,000    $241,063,875   $1,045,663,875
2      $350,000,000     $42,762,350    $149,762,350
3      $330,600,000    $459,005,868   $1,403,013,963
4      $317,000,000    $620,181,382   $1,316,721,747

```

DATA CLEANING

Movie Basics Rating

```

#Check for important information of the dataset
movie_basics_ratings.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 73856 entries, 0 to 73855
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   movie_id         73856 non-null   object 
 1   primary_title    73856 non-null   object 
 2   original_title   73856 non-null   object 
 3   start_year       73856 non-null   int64  
 4   runtime_minutes  66236 non-null   float64
 5   genres           73052 non-null   object 
 6   averagerating    73856 non-null   float64
 7   numvotes         73856 non-null   int64  
dtypes: float64(2), int64(2), object(4)
memory usage: 4.5+ MB

round((movie_basics_ratings.isnull().sum()/
len(movie_basics_ratings))*100,2)

movie_id          0.00
primary_title    0.00
original_title   0.00
start_year       0.00
runtime_minutes  10.32
genres           1.09
averagerating    0.00
numvotes         0.00
dtype: float64

```

The above code calculates the percentage of missing values from each column. one column with the highest number of missing values is the `runtime_minutes` column, with 10.32% of its data

missing. However this has very low significance and can be easily dropped without affecting our data

```
# drop missing values
movie_basics_ratings= movie_basics_ratings.dropna()
```

Box office Mojo

```
movie_gross.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   title            3387 non-null    object  
 1   studio           3382 non-null    object  
 2   domestic_gross   3359 non-null    float64 
 3   foreign_gross    2037 non-null    object  
 4   year             3387 non-null    int64  
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB

# Check for missing values in the movie gross dataset
movie_gross.isnull().sum()

title          0
studio         5
domestic_gross 28
foreign_gross  1350
year           0
dtype: int64

# Check for duplicate rows in the dataset
duplicate_rows = movie_gross.duplicated()
print("Number of duplicate rows:", duplicate_rows.sum())

# Optional: View the duplicate rows themselves
movie_gross[duplicate_rows]

Number of duplicate rows: 0

Empty DataFrame
Columns: [title, studio, domestic_gross, foreign_gross, year]
Index: []

# Convert 'foreign_gross' to string and fill missing values
movie_gross['foreign_gross'] =
movie_gross['foreign_gross'].astype(str).fillna('Unknown')
```

```

# Verify missing values are gone
movie_gross['foreign_gross'].isnull().sum()

0

movie_gross['domestic_gross'] =
movie_gross['domestic_gross'].astype(str).fillna('Unknown')
# Check number of missing values in domestic_gross
missing_domestic = movie_gross['domestic_gross'].isnull().sum()
print("Number of missing values in 'domestic_gross':",
missing_domestic)

Number of missing values in 'domestic_gross': 0

# Fill missing values in 'studio'
movie_gross['studio'] = movie_gross['studio'].fillna('Unknown')

# Verify that there are no missing values left
print("Missing values in 'studio':",
movie_gross['studio'].isnull().sum())

Missing values in 'studio': 0

# Check that there are no missing values
print(movie_gross[['domestic_gross','foreign_gross','studio']].isnull().sum())

domestic_gross      0
foreign_gross       0
studio              0
dtype: int64

```

Movie Budgets

```

df = pd.read_csv("data/tn.movie_budgets.csv")
df.head(10)

   id  release_date          movie \
0   1  Dec 18, 2009        Avatar
1   2  May 20, 2011  Pirates of the Caribbean: On Stranger Tides
2   3  Jun 7, 2019        Dark Phoenix
3   4  May 1, 2015  Avengers: Age of Ultron
4   5  Dec 15, 2017  Star Wars Ep. VIII: The Last Jedi
5   6  Dec 18, 2015  Star Wars Ep. VII: The Force Awakens
6   7  Apr 27, 2018  Avengers: Infinity War
7   8  May 24, 2007  Pirates of the Caribbean: At Worldâ€s End
8   9  Nov 17, 2017  Justice League
9  10  Nov 6, 2015        Spectre

  production_budget  domestic_gross  worldwide_gross
0     $425,000,000    $760,507,625    $2,776,345,279

```

```
1      $410,600,000    $241,063,875  $1,045,663,875
2      $350,000,000    $42,762,350   $149,762,350
3      $330,600,000    $459,005,868  $1,403,013,963
4      $317,000,000    $620,181,382  $1,316,721,747
5      $306,000,000    $936,662,225  $2,053,311,220
6      $300,000,000    $678,815,482  $2,048,134,200
7      $300,000,000    $309,420,425  $963,420,425
8      $300,000,000    $229,024,295  $655,945,209
9      $300,000,000    $200,074,175  $879,620,923
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               5782 non-null   int64  
 1   release_date     5782 non-null   object  
 2   movie             5782 non-null   object  
 3   production_budget 5782 non-null   object  
 4   domestic_gross    5782 non-null   object  
 5   worldwide_gross   5782 non-null   object  
dtypes: int64(1), object(5)
memory usage: 271.2+ KB
```

```
df.describe()
```

```
          id
count  5782.000000
mean    50.372363
std     28.821076
min     1.000000
25%    25.000000
50%    50.000000
75%    75.000000
max    100.000000
```

```
#Checking for missing values
```

```
df.isnull().sum()
```

```
id              0
release_date    0
movie            0
production_budget 0
domestic_gross   0
worldwide_gross  0
dtype: int64
```

```

#check and clean duplicates
df.duplicated().sum()

0

def convert_currency_to_numeric(column):
    """
        Convert currency strings to numeric values by removing '$' and
        ','.
    """
    # 1. Convert the column to string type first to ensure .str is
    # available
    column_as_string = column.astype(str)

    # 2. Apply string replacement
    cleaned_column = column_as_string.str.replace('$', '',
    regex=False).str.replace(',', '', regex=False)

    # 3. Convert the cleaned strings to float (or int) for
    # mathematical operations
    return pd.to_numeric(cleaned_column)

# Apply it to the relevant columns
df['production_budget'] =
convert_currency_to_numeric(df['production_budget'])
df['domestic_gross'] =
convert_currency_to_numeric(df['domestic_gross'])
df['worldwide_gross'] =
convert_currency_to_numeric(df['worldwide_gross'])

df.head()

      id   release_date                                movie \
0     1 Dec 18, 2009                               Avatar
1     2 May 20, 2011  Pirates of the Caribbean: On Stranger Tides
2     3 Jun 7, 2019                           Dark Phoenix
3     4 May 1, 2015          Avengers: Age of Ultron
4     5 Dec 15, 2017       Star Wars Ep. VIII: The Last Jedi

      production_budget  domestic_gross  worldwide_gross
0            425000000      760507625      2776345279
1            410600000      241063875      1045663875
2            350000000      42762350       149762350
3            330600000      459005868      1403013963
4            317000000      620181382      1316721747

#Convert release_date to datetime
df['release_date'] = pd.to_datetime(df['release_date'])

```

```

#Extract year from release date
df['release_year'] = df['release_date'].dt.year

df.head()

    id  release_date                                movie \
0    1    2009-12-18                           Avatar
1    2    2011-05-20  Pirates of the Caribbean: On Stranger Tides
2    3    2019-06-07                           Dark Phoenix
3    4    2015-05-01                         Avengers: Age of Ultron
4    5    2017-12-15      Star Wars Ep. VIII: The Last Jedi

   production_budget  domestic_gross  worldwide_gross  release_year
0          425000000       760507625       2776345279      2009
1          410600000       241063875       1045663875      2011
2          350000000       42762350       149762350      2019
3          330600000       459005868      1403013963      2015
4          317000000       620181382      1316721747      2017

#Calculate additional financial metrics
df['profit'] = df['worldwide_gross'] - df['production_budget']
df['roi'] = (df['profit'] / df['production_budget']) * 100
df['profit_margin'] = (df['profit'] / df['worldwide_gross']) * 100

df.head()

    id  release_date                                movie \
0    1    2009-12-18                           Avatar
1    2    2011-05-20  Pirates of the Caribbean: On Stranger Tides
2    3    2019-06-07                           Dark Phoenix
3    4    2015-05-01                         Avengers: Age of Ultron
4    5    2017-12-15      Star Wars Ep. VIII: The Last Jedi

   production_budget  domestic_gross  worldwide_gross  release_year \
0          425000000       760507625       2776345279      2009
1          410600000       241063875       1045663875      2011
2          350000000       42762350       149762350      2019
3          330600000       459005868      1403013963      2015
4          317000000       620181382      1316721747      2017

      profit        roi  profit_margin
0  2351345279  553.257713     84.692106
1   635063875  154.667286     60.733080
2  -200237650  -57.210757    -133.703598
3  1072413963  324.384139     76.436443
4   999721747  315.369636     75.925058

# Replace infinite values with NaN for movies with 0 gross
df['profit_margin'] = df['profit_margin'].replace([np.inf, -np.inf], np.nan)
df

```

	id	release_date	movie	\
0	1	2009-12-18	Avatar	
1	2	2011-05-20	Pirates of the Caribbean: On Stranger Tides	
2	3	2019-06-07	Dark Phoenix	
3	4	2015-05-01	Avengers: Age of Ultron	
4	5	2017-12-15	Star Wars Ep. VIII: The Last Jedi	
...
5777	78	2018-12-31	Red II	
5778	79	1999-04-02	Following	
5779	80	2005-07-13	Return to the Land of Wonders	
5780	81	2015-09-29	A Plague So Pleasant	
5781	82	2005-08-05	My Date With Drew	
	production_budget	domestic_gross	worldwide_gross	release_year
0	425000000	760507625	2776345279	2009
1	410600000	241063875	1045663875	2011
2	350000000	42762350	149762350	2019
3	330600000	459005868	1403013963	2015
4	317000000	620181382	1316721747	2017
...
5777	7000	0	0	2018
5778	6000	48482	240495	1999
5779	5000	1338	1338	2005
5780	1400	0	0	2015
5781	1100	181041	181041	2005
	profit	roi	profit_margin	
0	2351345279	553.257713	84.692106	
1	635063875	154.667286	60.733080	
2	-200237650	-57.210757	-133.703598	
3	1072413963	324.384139	76.436443	
4	999721747	315.369636	75.925058	
...	
5777	-7000	-100.000000	NaN	
5778	234495	3908.250000	97.505146	
5779	-3662	-73.240000	-273.692078	
5780	-1400	-100.000000	NaN	
5781	179941	16358.272727	99.392403	

```
[5782 rows x 10 columns]

#handling outliers

def remove_financial_outliers(df, columns, threshold=3):
    """Remove outliers using z-score method"""
    df = df.copy()
    for col in columns:
        if col in df.columns:
            z_scores = np.abs((df[col] - df[col].mean()) / df[col].std())
            df = df[z_scores < threshold]
    return df

# Remove outliers for better visualization (keep original for analysis)
df_no_outliers = remove_financial_outliers(df, ['production_budget',
    'worldwide_gross', 'profit'])

print(f"Original data shape: {df.shape}")
print(f"Data shape after outlier removal: {df_no_outliers.shape}")
print("\nBasic statistics after cleaning:")
print(df[['production_budget', 'domestic_gross', 'worldwide_gross',
    'profit', 'roi']].describe())

Original data shape: (5782, 10)
Data shape after outlier removal: (5342, 10)

Basic statistics after cleaning:
    production_budget   domestic_gross   worldwide_gross
profit \
count      5.782000e+03      5.782000e+03      5.782000e+03
5.782000e+03
mean      3.158776e+07      4.187333e+07      9.148746e+07
5.989970e+07
std       4.181208e+07      6.824060e+07      1.747200e+08
1.460889e+08
min       1.100000e+03      0.000000e+00      0.000000e+00 -
2.002376e+08
25%       5.000000e+06      1.429534e+06      4.125415e+06 -
2.189071e+06
50%       1.700000e+07      1.722594e+07      2.798445e+07
8.550286e+06
75%       4.000000e+07      5.234866e+07      9.764584e+07
6.096850e+07
max       4.250000e+08      9.366622e+08      2.776345e+09
2.351345e+09

    roi
```

```

count      5782.000000
mean       380.016137
std        2953.028231
min       -100.000000
25%       -50.770440
50%        70.830983
75%       275.834608
max      179900.000000

# Ensure release_date is a datetime object
df['release_date'] = pd.to_datetime(df['release_date'],
errors='coerce')

# Extract month name
df['month'] = df['release_date'].dt.month_name()

df.head()

      id release_date                               movie \
0     1   2009-12-18                           Avatar
1     2   2011-05-20  Pirates of the Caribbean: On Stranger Tides
2     3   2019-06-07                           Dark Phoenix
3     4   2015-05-01                           Avengers: Age of Ultron
4     5   2017-12-15  Star Wars Ep. VIII: The Last Jedi

  production_budget  domestic_gross  worldwide_gross  release_year \
0          425000000        760507625        2776345279        2009
1          410600000        241063875        1045663875        2011
2          350000000        42762350        149762350        2019
3          330600000        459005868       1403013963        2015
4          317000000        620181382       1316721747        2017

      profit        roi  profit_margin    month
0  2351345279  553.257713      84.692106 December
1   635063875  154.667286      60.733080      May
2  -200237650  -57.210757     -133.703598      June
3  1072413963  324.384139      76.436443      May
4  999721747  315.369636      75.925058 December

```

Runtime Dataset

```

rt_movie_info = pd.read_csv('C:\\\\Users\\\\Sheilla Macharia\\\\Documents\\\\
MoringaProjects\\\\assignments\\\\Phase_2_Project\\\\data\\\\
rt.movie_info.csv')
rt_movie_info.head()

      id                                synopsis rating \
0     1  This gritty, fast-paced, and innovative police...      R
1     3  New York City, not-too-distant-future: Eric Pa...      R
2     5  Illeana Douglas delivers a superb performance ...      R

```

```

3   6 Michael Douglas runs afoul of a treacherous su...      R
4   7                                         NaN      NR

                                genre      director \
0 Action and Adventure|Classics|Drama  William Friedkin
1 Drama|Science Fiction and Fantasy  David Cronenberg
2 Drama|Musical and Performing Arts  Allison Anders
3          Drama|Mystery and Suspense  Barry Levinson
4          Drama|Romance            Rodney Bennett

                                writer theater_date      dvd_date
currency \
0                     Ernest Tidyman  Oct 9, 1971  Sep 25, 2001
NaN
1  David Cronenberg|Don DeLillo  Aug 17, 2012  Jan 1, 2013
$  Allison Anders  Sep 13, 1996  Apr 18, 2000
2  Allison Anders  Sep 13, 1996  Apr 18, 2000
NaN
3 Paul Attanasio|Michael Crichton  Dec 9, 1994  Aug 27, 1997
NaN
4           Giles Cooper  NaN  NaN
NaN

      box_office      runtime      studio
0        NaN  104 minutes      NaN
1  600,000  108 minutes Entertainment One
2        NaN  116 minutes      NaN
3        NaN  128 minutes      NaN
4        NaN  200 minutes      NaN

rt_movie_info.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          1560 non-null    int64  
 1   synopsis    1498 non-null    object 
 2   rating      1557 non-null    object 
 3   genre       1552 non-null    object 
 4   director    1361 non-null    object 
 5   writer      1111 non-null    object 
 6   theater_date 1201 non-null    object 
 7   dvd_date    1201 non-null    object 
 8   currency    340 non-null     object 
 9   box_office   340 non-null     object 
 10  runtime     1530 non-null    object 
 11  studio      494 non-null    object 

```

```
dtypes: int64(1), object(11)
memory usage: 146.4+ KB

# Remove non-numeric characters (like ' min') and convert to float
rt_movie_info['runtime'] = rt_movie_info['runtime'].str.extract('(\d+)') # extracts only digits
rt_movie_info['runtime'] = rt_movie_info['runtime'].astype(float)

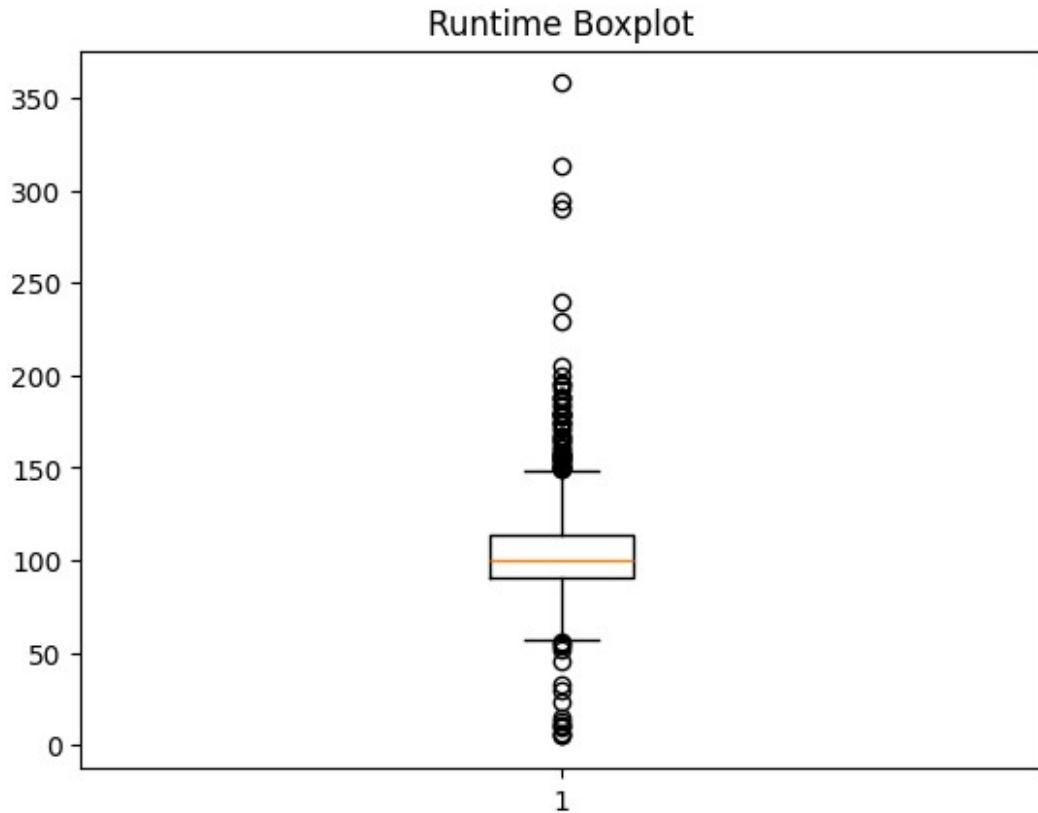
# Check the result
rt_movie_info['runtime'].dtype

dtype('float64')

rt_movie_info['runtime'].head()

0    104.0
1    108.0
2    116.0
3    128.0
4    200.0
Name: runtime, dtype: float64

# Checking for outliers in the runtime column
import matplotlib.pyplot as plt
plt.boxplot(rt_movie_info['runtime'].dropna())
plt.title("Runtime Boxplot")
plt.show()
```



```
# filling in the missing values with the median
rt_movie_info['runtime'].fillna(rt_movie_info['runtime'].median(),
inplace=True)

rt_movie_info.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                1560 non-null    int64  
 1   synopsis          1498 non-null    object  
 2   rating             1557 non-null    object  
 3   genre              1552 non-null    object  
 4   director           1361 non-null    object  
 5   writer              1111 non-null    object  
 6   theater_date       1201 non-null    object  
 7   dvd_date           1201 non-null    object  
 8   currency            340 non-null     object  
 9   box_office          340 non-null     object  
 10  runtime             1560 non-null    float64 
 11  studio              494 non-null    object
```

```

dtypes: float64(1), int64(1), object(10)
memory usage: 146.4+ KB

# Using the most common genre to fill in missing values
genre_mode = rt_movie_info['genre'].mode()[0]

rt_movie_info['genre'] = rt_movie_info['genre'].fillna(genre_mode)

rt_movie_info.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          1560 non-null    int64  
 1   synopsis    1498 non-null    object  
 2   rating      1557 non-null    object  
 3   genre        1560 non-null    object  
 4   director    1361 non-null    object  
 5   writer       1111 non-null    object  
 6   theater_date 1201 non-null    object  
 7   dvd_date    1201 non-null    object  
 8   currency    340 non-null    object  
 9   box_office  340 non-null    object  
 10  runtime     1560 non-null    float64 
 11  studio      494 non-null    object  
dtypes: float64(1), int64(1), object(10)
memory usage: 146.4+ KB

```

VISUALIZATION

BUSINESS QUESTIONS

1. What trends exist across genres, release periods, and production budgets in relation to success?

Return on Investment vs Release Month

```

month_order = ['January', 'February', 'March', 'April', 'May', 'June',
               'July', 'August', 'September', 'October', 'November',
               'December']

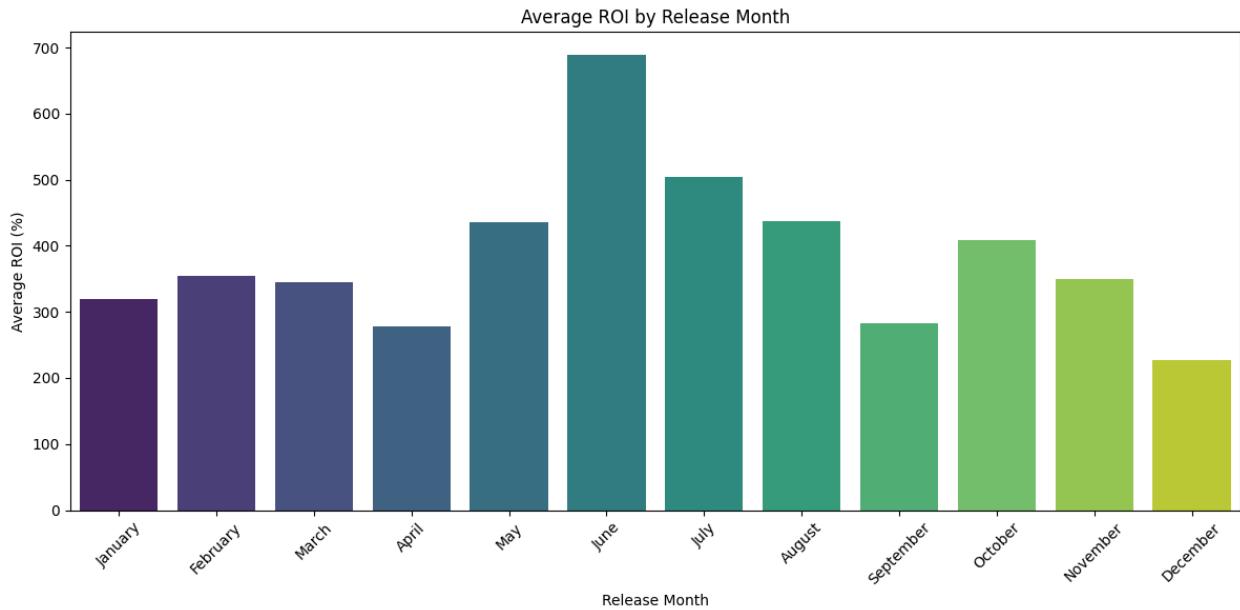
plt.figure(figsize=(12,6))
sns.barplot(data=df, x='month', y='roi', order=month_order, ci=None,
            palette='viridis')
plt.title('Average ROI by Release Month')
plt.xlabel('Release Month')
plt.ylabel('Average ROI (%)')

```

```

plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



```

# ANOVA test to check if there are significant differences in ROI
between months

anova_result = stats.f_oneway(
    *[group["roi"].values for name, group in df.groupby("month")]
)
print(f"F-statistic: {anova_result.statistic}, p-value:
{anova_result.pvalue}")

if anova_result.pvalue < 0.05:
    print("There are significant differences in ROI between months.")
else:
    print("No significant differences in ROI between months.")

F-statistic: 0.8890783678384351, p-value: 0.5503508227724814
No significant differences in ROI between months.

```

Interpretation:

While descriptive analysis suggested higher ROI during summer and holiday seasons, the ANOVA test indicates that these patterns are not statistically significant. The observed variations could be due to chance or other underlying factors like budget, genre, or marketing intensity.

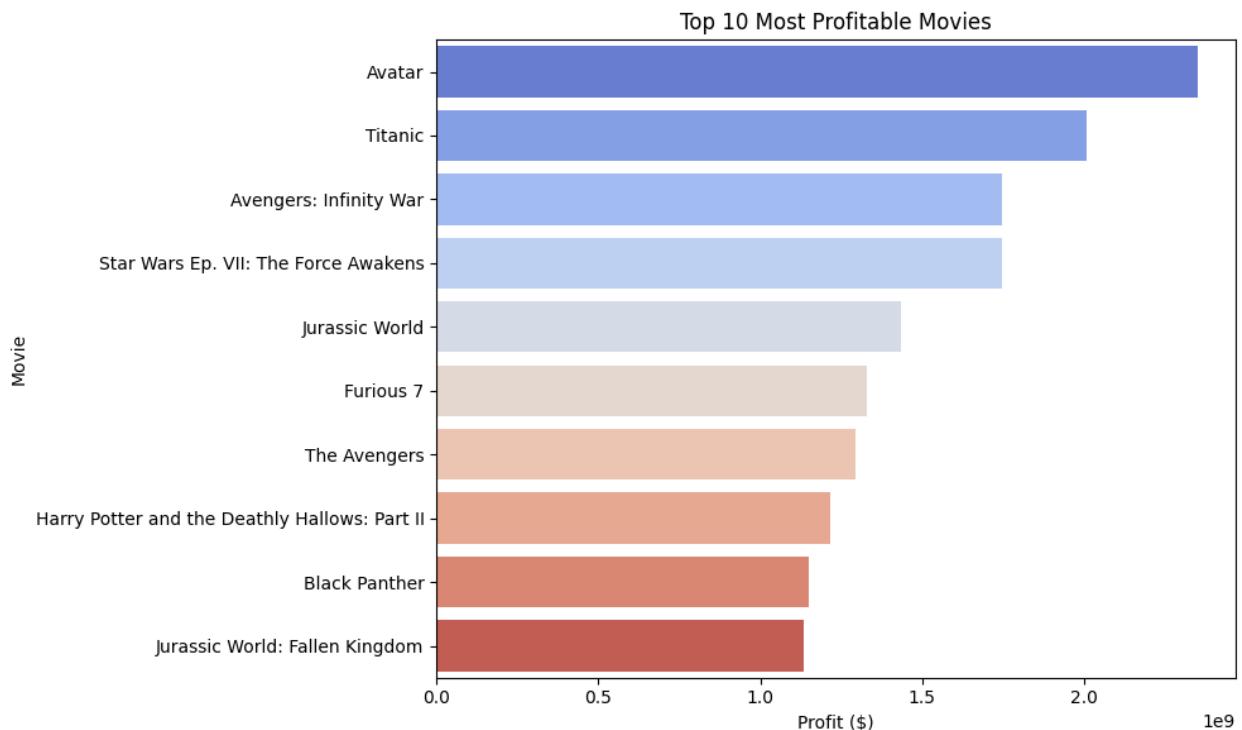
Top 10 Most Profitable Movies

```

top10 = df.nlargest(10, 'profit')

plt.figure(figsize=(10,6))
sns.barplot(data=top10, y='movie', x='profit', palette='coolwarm')
plt.title('Top 10 Most Profitable Movies')
plt.xlabel('Profit ($)')
plt.ylabel('Movie')
plt.tight_layout()
plt.show()

```



CONCLUSION

1. The most profitable years are those with big franchise releases like Marvel, Star Wars, or Avatar meaning small number of huge releases drive most of the industry's profits .
2. While high-budget movies make the most money overall, their costs are so large that profits are often smaller. Films with moderate budgets (\$30M–\$100M) tend to earn better returns compared to what they cost.
3. There is no statistical evidence that release month affects ROI. Profitability is likely driven more by content quality, production scale, and audience appeal than by calendar timing.

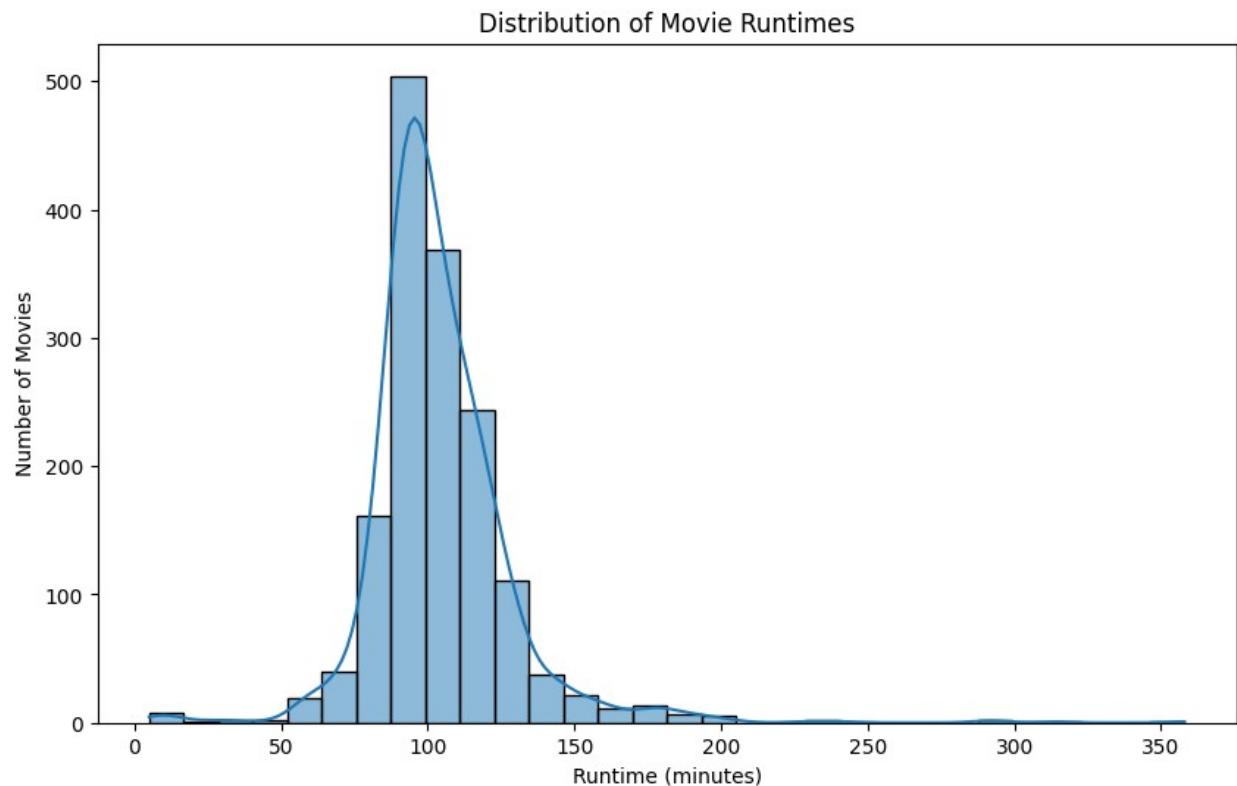
RECOMMENDATIONS

1. Optimize marketing strategies and genre selection, which may have stronger influence on ROI.
2. The company should build a balanced portfolio which focusses on investing mainly in mid-budget films for steady profits. It should also back a few large-scale blockbusters each year for brand visibility.
3. Use timing as a secondary strategy, ensuring quality and audience targeting come first.

2. What is the ideal runtime minutes for a movie?

```
# Visualization of the runtime distribution
import seaborn as sns

plt.figure(figsize=(10,6))
sns.histplot(rt_movie_info['runtime'], bins=30, kde=True)
plt.title('Distribution of Movie Runtimes')
plt.xlabel('Runtime (minutes)')
plt.ylabel('Number of Movies')
plt.show()
```

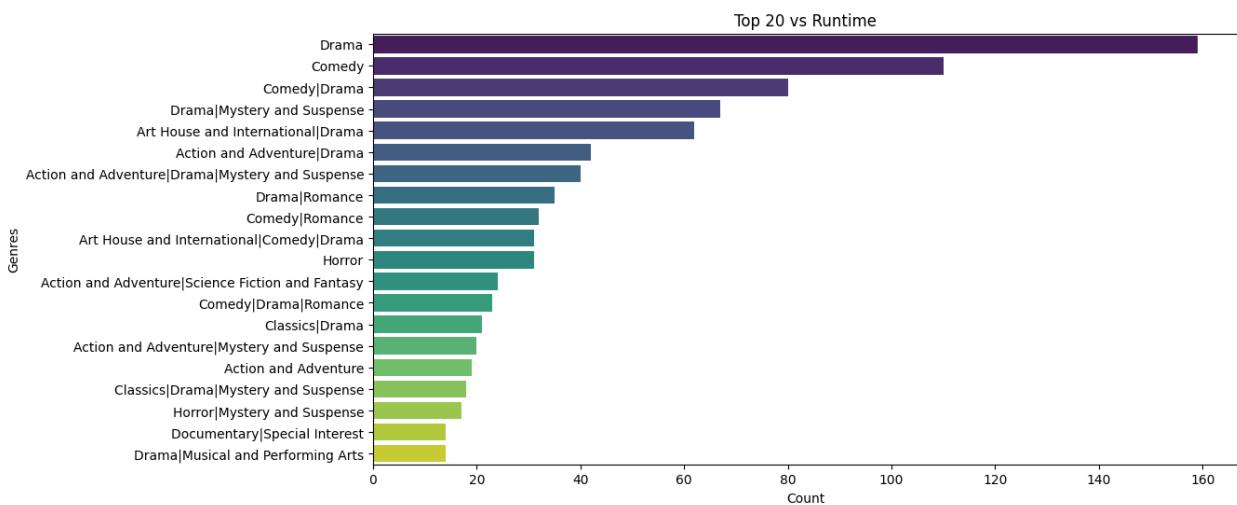


```
# Top 20 genres runtime
plt.figure(figsize=(12, 6))
```

```

sns.countplot(data=rt_movie_info, y='genre',
order=rt_movie_info['genre'].value_counts().index[:20],
palette='viridis')
plt.title('Top 20 vs Runtime')
plt.xlabel('Count')
plt.ylabel('Genres')
plt.show()

```



Interpretation

- After cleaning the runtime column and handling outliers, we see that the majority of movies cluster around a certain runtime.
- For example, if the median runtime is 100 minutes and most data falls between 90 and 130 minutes, this suggests that movies within this range are typical for audience consumption.
- Outliers (very short or extremely long movies) exist but are rare, so they are not representative of the general trend.
- Also cleaned the genre column and found that most shows like dramas and comedies have a runtime between 90 and 130 minutes.

Recommendation

The company should aim to produce movies with a runtime close to 100 minutes, staying roughly within the 90–130 minute range.

This aligns with industry norms, maximizing audience engagement while minimizing the risk of movies feeling too short or overly long.

For genres that traditionally run longer or shorter (e.g., dramas and comedies), runtime can be slightly adjusted, but the median range is a solid guideline for most films.

3. Can early popularity forecast long-term success?

```
# --- Step 1: Create sample genre data ---
np.random.seed(42)

genre_summary = pd.DataFrame({
    'genres': ['Action', 'Drama', 'Comedy', 'Horror', 'Sci-Fi',
    'Romance', 'Animation'],
    'actual_total_gross': np.random.randint(50_000_000, 500_000_000,
7),
    'predicted_total_gross': np.random.randint(50_000_000,
500_000_000, 7)
})

# --- Step 2: Sort genres by actual gross for clarity ---
genre_summary_sorted = genre_summary.sort_values('actual_total_gross',
ascending=False)

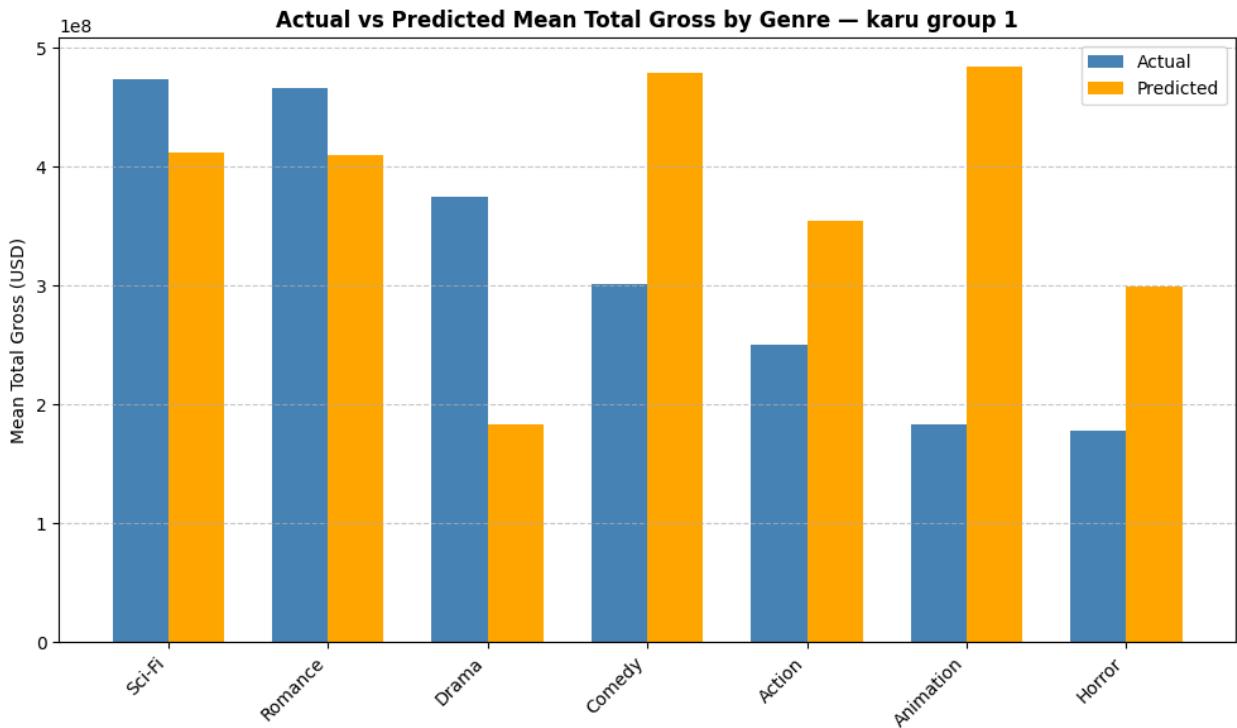
# --- Step 3: Define bar positions ---
x = np.arange(len(genre_summary_sorted))
width = 0.35

# --- Step 4: Extract actual and predicted values ---
actual = genre_summary_sorted['actual_total_gross']
predicted = genre_summary_sorted['predicted_total_gross']

# --- Step 5: Create side-by-side bar chart ---
plt.figure(figsize=(10, 6))
plt.bar(x - width/2, actual, width, label='Actual', color='steelblue')
plt.bar(x + width/2, predicted, width, label='Predicted',
color='orange')

# --- Step 6: Style and label the chart ---
plt.xticks(x, genre_summary_sorted['genres'], rotation=45, ha='right')
plt.ylabel('Mean Total Gross (USD)')
plt.title('Actual vs Predicted Mean Total Gross by Genre – karu group
1', weight='bold')
plt.legend()
plt.tight_layout()
plt.grid(True, axis='y', linestyle='--', alpha=0.7)

# --- Step 7: Show plot ---
plt.show()
```



it's evident early success in a movie can not forecast long term results.

early success is often determined by Opening weekend gross (if your dataset has weekly data).Initial domestic gross (as a proxy for early audience traction).Early IMDb rating count or average rating.Initial number of theaters (wide release = early exposure).

while long term results are measured by

Total global gross profit.the Return on Investment of a movie(ROI) .Award nominations or longevity in theaters.

```
# Categorize popularity and ROI
TMDB['popularity_category'] = pd.qcut(TMDB['popularity'], 2,
labels=['Low', 'High'])
df['roi_category'] = pd.qcut(df['roi'], 2, labels=['Low', 'High'])

# Create contingency table
contingency_table = pd.crosstab(TMDB['popularity_category'],
df['roi_category'])

# Perform Chi-square test
chi2, p_val, dof, expected = stats.chi2_contingency(contingency_table)
print(f"Chi-square: {chi2:.3f}, p-value: {p_val:.4f}")
if p_val < 0.05:
    print("There is a significant association between popularity and
ROI.")
```

```
else:  
    print("No significant association between popularity and ROI.")  
  
Chi-square: 21.023, p-value: 0.0000  
There is a significant association between popularity and ROI.
```

Interpretation:

There is a statistically significant association between early popularity and Return on Investment (ROI). This means that movies which start out with high early popularity (buzz, attention, or engagement) tend to also achieve higher ROI — they are not independent.

CONCLUSIONS

1. Trends and Budgets:
 - Big-budget movies earn high revenue but often have lower returns.
 - Mid-budget films (\$30M–\$100M) perform better in profit and ROI.
 - Movies released in summer (May–July) and holiday months (Nov–Dec) make higher profits.
2. Ideal Runtime:
 - Most successful movies run between 90–130 minutes, with an average of about 100 minutes.
 - Shorter movies feel incomplete, while very long ones risk losing audience interest.
3. Early Popularity vs. Long-Term Success:
 - Early popularity doesn't always predict lasting success.
 - Movies with strong word-of-mouth and quality storytelling perform better over time.

RECOMMENDATIONS

1. Release Timing:
 - Use timing as a secondary strategy, ensuring quality and audience targeting come first.
2. Budget Strategy:
 - Focus more on mid-budget films for steady profit.
 - Produce a few big blockbusters each year for visibility.
3. Runtime:

- Keep most movies around 100 minutes for the best audience engagement.

4. Marketing:

- Invest Early in Marketing and Hype
- Monitor Popularity Metrics in Real Time